
Calibrated Risk Envelopes for LLM Quantization, and the Limits of Predicting Compression Damage

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Quantizing a large language model reduces serving cost and costs some accu-
2 racy. Determining how much normally requires re-running a full benchmark suite,
3 which is the expense quantization was meant to avoid. We ask whether the accu-
4 racy delta can be predicted instead. Using 817 published evaluations spanning 38
5 base checkpoints, 8 model families and 6 quantization schemes, we report three
6 results. First, a *negative* one: per-model prediction carries almost no signal. Un-
7 der leave-one-family-out validation the best predictor, a per-scheme mean, beats
8 a global mean by 0.03pp of MAE (95% checkpoint-bootstrap CI [0.01, 0.05]),
9 and both ridge regression and gradient boosting are *worse* than guessing the av-
10 erage, because evaluation noise accounts for at least a third of the variance on 8
11 of 11 benchmarks. Second, a positive one: conditioning only on the quantization
12 scheme yields split-conformal intervals with 90.1% empirical coverage (118/131,
13 95% CI [83.6, 94.6]) on checkpoints never used in construction. Those check-
14 points cover 4 of the 6 schemes, and W4A16 supplies 80 of the 131 rows. Each
15 (scheme, size-band) cell returns one of three verdicts: trusted; insufficient evi-
16 dence, where the interval is still shown but the scheme is demoted and flagged;
17 or refused, where the interval is withheld. Currently 9 of 17 cells are insufficient
18 evidence and 0 are refused. Third, and most consequential for practice: the pub-
19 lished corpus is outcome-truncated. Every configuration in it is one that someone
20 chose to release. We rented an A100 and measured deliberately-faulted configu-
21 rations, observing losses to -39.5pp against a published worst case of -8.86pp —
22 measured, we stress, only on models of 1.5B and below, and on configurations
23 engineered to fail rather than typical ones. This is an existence proof that far
24 worse outcomes exist than the corpus contains; it does not estimate how common
25 they are. We argue the resulting artifact should be described as a calibrated his-
26 torical baseline rather than a predictor, and we state precisely why its coverage
27 guarantee does not transfer to a practitioner’s own untuned recipe. Code, data and
28 the provenance of every reported figure are provided as supplementary material.

29 1 Introduction

30 Post-training quantization is now routine. A practitioner choosing between W4A16, W8A8 and FP8
31 wants one number: how much accuracy will this cost? In practice the only way to know is to run the
32 evaluation.

33 This paper asks whether that number can be forecast from the public record. Model publishers
34 release quantized checkpoints with benchmark tables comparing the compressed model against its
35 dense parent. Aggregated, these are hundreds of paired measurements of exactly the quantity of
36 interest (850 extracted in our harvest).

37 Our contributions:

- 38 1. **A negative result** (§5). Per-model prediction of the accuracy delta does not work, and
39 we show why: the target is dominated by evaluation noise. We report this as the primary
40 finding rather than burying it.
- 41 2. **A calibrated envelope that does work** (§6), conditioned on scheme alone, with measured
42 coverage on held-out families and a three-state verdict for each (scheme, size-band) cell:
43 *trusted*; *insufficient evidence*, where the interval is still shown but the scheme is demoted
44 and the cell flagged as one the tool cannot judge; or *refused*, where the interval is withheld.
45 Currently 9 of 17 cells are insufficient evidence and 0 are refused.
- 46 3. **An existence proof that the corpus omits severe outcomes** (§7), via GPU experiments on
47 deliberately-faulted configurations—data the published record structurally cannot contain.
48 It shows that losses far beyond anything published are reachable on small models. It does
49 not estimate how large the bias is: our attempt at a correction (in the supplement) did not
50 work well enough to ship.
- 51 4. **An audit** (§9) recording errors we found in our own work, including ones that changed
52 published numbers and claims.

53 2 Related Work

54 **Predicting evaluations.** BenchPress Zeng and Papailiopoulos [2026] establishes the mechanism
55 closest to ours: predict an eval result, estimate whether to trust it, wrap it in a conformal interval, and
56 decide whether to skip the real run. Their score matrix over 84 frontier models and 133 benchmarks
57 is approximately rank-2. We do not claim novelty against them on mechanism. Their estimand is
58 a model’s absolute score; ours is the paired delta between a model and its compressed self, and
59 quantized checkpoints do not appear in their paper.

60 **Conformal prediction on compressed models.** Tong et al. [2026] apply conformal prediction
61 directly to quantized and sparse LLMs, which is a closer domain than BenchPress. The estimand
62 differs: their conformal sets are over label space, constructed from a compressed model’s own output
63 probabilities, so the method measures uncertainty *after* running the model. Ours is constructed over
64 historical accuracy deltas to give an interval *before* running anything. Their Figure 2 plots the
65 same quantity we predict, $\text{Acc}_{\text{compressed}} - \text{Acc}_{\text{dense}}$, but measures it. Their finding that compression
66 decouples accuracy from uncertainty bounds what an accuracy-only estimate like ours can claim
67 (§8).

68 **Quantization methods.** The checkpoints we harvest are produced with GPTQ Frantar et al. [2022]
69 and SmoothQuant Xiao et al. [2022] via llm-compressor vLLM Project [2024]. Kumar et al. [2024]
70 fit a functional form that predicts degradation from training and inference in varied precisions, using
71 pretraining runs of models up to 1.7B parameters. We fit no such law: our envelope is empirical, per
72 scheme, over published post-training-quantized checkpoints.

73 **The source of our data.** Kurtić et al. [2024] report over 500,000 evaluations of quantized Llama-
74 3.1 models, finding recovery above 99% of the average OpenLLM v1 score for every scheme tested.
75 As far as we can tell, that work does not address whether the published set of quantized models is a
76 biased sample. Characterising that gap is the contribution of §7.

77 **Selection bias.** Survivorship bias has been identified and measured in a widely used retrieval
78 benchmark Gupta and MacAvaney [2022]. Conformal prediction’s guarantee is known to fail un-
79 der selection: conditional on a selection event, calibration data are no longer exchangeable with a
80 test point Barber et al. [2023], Jin and Ren [2024]. We apply this directly to our own artifact in §8.

81 3 Data

82 We harvest published model cards into rows of (model, scheme, benchmark, $\text{acc}_{\text{before}}$, $\text{acc}_{\text{after}}$), yield-
83 ing 850 rows, of which 817 remain after dropping 33 whose baseline accuracy is below 20% (near
84 chance, where a delta is uninformative), across 38 base checkpoints, 8 families and 6 schemes.

85 **An integrity gate.** Cards typically print a *Recovery %* column alongside the before/after pair.
 86 This is redundant information, and we exploit it: for every three-column row we check recovery $\approx$
 87 $100 \times \text{after/before}$ to a tolerance derived from the card’s own printed precision, and reject rows that
 88 disagree. This caught spurious values that our parser produced under column-order variation be-
 89 tween card templates, and internally inconsistent cards. Rejections are recorded rather than silently
 90 dropped.

91 **A structural limitation, stated at the outset.** All rows come from a single publisher using a
 92 single toolchain. Holding out a model family tests transfer across *models*; nothing in this design
 93 tests transfer across *quantization practitioners*. §7 gives direct evidence that this distinction matters.

94 4 Method

95 **Split conformal.** Following split conformal regression Lei et al. [2018], given calibration residuals
 96 $\{s_i\}_{i=1}^n$ and level α , the interval half-width is the $\lceil (n+1)(1-\alpha) \rceil$ -th smallest score, which gives
 97 exact finite-sample coverage under exchangeability. Where n is too small for that index to exist, we
 98 return ∞ rather than a misleadingly finite band.

99 **Mondrian (group-conditional) calibration.** Calibrating separately within categories is the Mon-
 100 drian construction of Vovk et al. [2005], applied to regression by Boström and Johansson [2020]. A
 101 single marginal interval advertised at 90% was in fact a 68% interval for NVFP4 users—the users
 102 who most need it. We therefore calibrate per scheme.

103 **A three-state verdict.** For each (scheme, size-band) cell we measure coverage directly, and the
 104 tool returns one of three verdicts. A cell is *trusted* when measured coverage is adequate and ad-
 105 equately supported. It is *refused*—INSUFFICIENT CALIBRATION, no interval—when coverage is
 106 measurably poor: the checkpoint-bootstrap interval on its coverage lies wholly below the 85% line,
 107 on at least 3 checkpoints. It is *insufficient evidence* when the cell rests on fewer than 3 distinct
 108 checkpoints, or when a checkpoint-level cluster bootstrap 90% interval on its coverage straddles the
 109 85% line. In that case the evidence cannot support either of the other two verdicts. The tool still
 110 prints the all-sizes interval for such a cell, but demotes the scheme to the lowest tier and states that
 111 the cell cannot be judged; only a refused cell has its interval withheld.

112 Coverage is scored *one-sided*: how often the true delta stayed at or above the interval’s lower bound.
 113 The risk this tool bounds is accuracy loss, and a model that beats its envelope has exposed no one
 114 to anything, so counting that as a failure would flag cells that are merely outperforming. Comparing
 115 a one-sided measurement against the printed 90% is the *lenient* reading, not a flattering one: the
 116 interval is symmetric, so the permitted 10% is split across two tails, and pooled one-sided coverage,
 117 across all 817 rows with each row scored under 7 different calibration families, is 95.1%.

118 Currently 0 of 17 cells are refused. Two cells were flagged under the earlier two-state rule, and nei-
 119 ther is a confident refusal under the three-state one. `w4a16|<2B` had one-sided coverage of 70.1%
 120 over 11 rows (two-sided 68.8%), but those rows come from only 2 checkpoints. `w8a16|>10B` had
 121 81.1% over 31 rows from 5 checkpoints (two-sided 73.7%), and its bootstrap interval, [47, 100]%,
 122 straddles the threshold. Both are now *insufficient evidence*. We had scoped the change to those
 123 two cells; applying the rule to every cell moved 9 of 17 into that state, including several whose
 124 point-estimate coverage looked comfortable. The same measurement and the same rule judge each
 125 scheme’s all-sizes coverage: 1 of the 6 schemes is insufficient evidence at that level and 0 is refused.
 126 The one-sided scoring and the three-state rule were proposed by an external commenter who identi-
 127 fies as an AI account collaborating with the OpenLanguageModel maintainers; we recomputed each
 128 of their figures from the raw rows before adopting either change, and one of their descriptions did
 129 not survive that check.

130 **A support floor.** Coverage measured on many benchmarks from one checkpoint is not indepen-
 131 dent evidence; it is one quantization run viewed many ways. Cells resting on fewer than 3 distinct
 132 checkpoints cannot reach the top tier regardless of measured coverage. 4 cells are affected.

5 The negative result

Under leave-one-family-out validation, row-weighted MAE:

predictor	MAE (pp)	beats global mean?
per-scheme mean (shipped)	0.7227	yes
global mean (baseline)	0.7545	—
per-(scheme $\times$ benchmark)	0.7607	no
ridge regression	0.7639	no
per-benchmark mean	0.7647	no
gradient boosting	0.7878	no

The per-scheme mean’s advantage is small: 0.03pp, with a checkpoint-bootstrap 95% CI of [0.01, 0.05]pp, and it is positive in 7 of the 8 held-out families. It is a distinguishable but modest gain. Ridge and gradient boosting are both *worse* than predicting the global average. Adding model size, bit-width, method, benchmark identity and headroom each degraded out-of-family performance.

Why. Estimating per-benchmark evaluation noise from near-lossless schemes—whose true delta should be approximately zero—gives a noise floor that accounts for anywhere from 15% of the observed variance to all of it across 11 benchmarks. On GSM8K the deltas of supposedly lossless schemes have a standard deviation of 2.1pp, while the quantity being predicted has mean -0.39 pp. The target is a difference of two noisy quantities that largely cancel: the numerically smallest and hardest-to-predict summary of quantization damage available. Two cautions on the noise figures. They assume near-lossless schemes have a true delta near zero, so they are upper bounds on noise. And 5 of the 11 estimates rest on 10 or fewer rows; one, at 124%, exceeds the total variance, which we attribute to sampling error rather than treat as a finding.

Implication. The entire fitted artifact is 18 numbers. A raw per-scheme historical quantile band, with no model and no conformal machinery, is a serious competitor: it looked better covered in-sample, though under the strict protocol that advantage disappears (§9). What conformal contributes is a finite-sample guarantee and a principled rule for small- n schemes, not accuracy. We therefore describe the result as a *calibrated historical baseline*, not a predictor.

6 Results

scheme	90% interval (pp)	worst observed	% losing >3pp	evals
fp8	[-1.11, +0.48]	-1.90pp	0.0%	80
w8a8_int	[-1.71, +1.07]	-3.32pp	1.0%	194
fp8_dynamic	[-1.44, +1.34]	-8.72pp	1.0%	200
w8a16	[-0.98, +0.81]	-3.04pp	1.2%	86
w4a16	[-2.87, +1.40]	-8.86pp	5.7%	193
nvfp4	[-4.10, +1.87]	-8.12pp	15.6%	64

Prospective validation. After freezing the method we harvested checkpoints not used during development. Under a strict protocol that excludes any family present in training and any card whose format informed a parser fix, coverage is $118/131 = 90.1\%$, 95% Clopper–Pearson CI [83.6, 94.6], against a nominal 90%. A from-scratch, standard-library-only reimplementation sharing no code with the pipeline reproduces this figure with zero disagreements. That checks the computation, not the validity of the design.

What the prospective set does and does not cover. The 131 strict rows come from 19 quantized checkpoints, not 131 independent observations: a bootstrap that resamples checkpoints gives a 95% interval of [83.2, 94.7]%, close to the row-level one. The rows cover only 4 of the 6 schemes; FP8 and NVFP4 have no prospective rows, so the figure says nothing about them.

scheme	strict rows	checkpoints	covered (%)
w4a16	80	11	87.5
w8a8_int	36	5	91.7
fp8_dynamic	11	2	100.0
w8a16	4	1	100.0

W4A16, the most common scheme, dominates the pooled figure and sits below the nominal 90%; the other schemes rest on few checkpoints. Over all 186 harvested prospective rows, before the strict filter, coverage is 90.3%, so the filter did not manufacture the headline.

What the envelope cannot do. 0 of 6 intervals exclude zero: the tool can never state that a scheme *will* cost accuracy. And 19 of 25 observed losses worse than 3pp fell *below* their interval floor—when real damage occurs, the interval usually misses it. This is the single most important caveat in the paper.

7 The measured left tail

The corpus contains only configurations someone chose to publish. To measure what it structurally cannot show, we rented an A100 and quantized real models under deliberately-faulted recipes.

Result. Worst measured loss -39.5pp , against -8.86pp as the worst loss anywhere in the published corpus. 16 of 36 faulted rows lost 3pp or more; separately, 2 of 9 rows from our correctly-configured control arm did too. The dominant single effect was wrong scale granularity—per-tensor instead of per-group.

Scope, stated plainly. This tail was measured only on models of 1.5B and below. It establishes that catastrophic damage is reachable; it is not a measured bound for larger models.

A confound we found and did not resolve. Re-running MMLU under the published protocol reproduces published base accuracy (47.20 against 47.42; 59.57 against 60.98), closing the harness question. It exposed a larger one: our own *correctly-configured* control is $2.6\times$ more damaging than published W4A16 on closely comparable checkpoints (mean gap -3.78pp , $z = 2.9$ and 5.3). We attribute this to approximately $171\times$ less calibration data than production defaults; we found no bug. Consequently, excess-over-control measures *bad recipe versus mediocre recipe*, not versus production. It is a conservative anchor, not a bound, and we decline to publish a corrected point estimate: the correction moves by 5.70pp depending on the anchor set, and we have only 6 rows from two models.

8 Limitations

The guarantee is conditional, and absent where it matters most. Split conformal gives exact coverage under exchangeability between calibration and test. Our calibration set is the set of configurations a publisher chose to release—a selection event. Conditional on selection, calibration points are not exchangeable with an arbitrary test point and the guarantee does not transfer Barber et al. [2023], Jin and Ren [2024]. It holds for a new checkpoint drawn from the same publication process. It carries **no guarantee** for a recipe the reader tuned themselves and that no publisher would have released—which is the case a practitioner most wants covered.

Accuracy is not the whole of deployment risk. Tong et al. [2026] show compression frequently decouples accuracy from uncertainty: a model can hold accuracy while its prediction sets inflate. A scheme this work calls low-risk on accuracy may still degrade uncertainty quality. We measure none of that.

One publisher, one toolchain, and why. See §3. Our own control arm (§7) is direct evidence that practitioner and calibration volume can dominate scheme choice. We measured how confining this is: surveying 10 publishers and 260 quantized model cards, 10 of 222 cards from the other 9

207 publishers carry paired before/after evaluations, 9 of them Intel’s, so others do publish them. What
 208 the other publishers almost never provide is a *Recovery* column, which is what lets each row be
 209 checked against $100 \times$ after/before and rejected on disagreement: of their rows, 6 pass that gate,
 210 all from 1 card, against 147 from the control publisher. The corpus is single-publisher not because
 211 others publish nothing, but because the control publisher’s cards are, by a wide margin, the ones
 212 published in a self-verifying form.

213 **What can and cannot be re-run.** The supplement recomputes every tagged figure from the data
 214 files it ships. It cannot re-run the faulted-configuration experiments (they need a GPU; the recorded
 215 results are included), the harvest from raw model cards (the cards are not redistributed), or the
 216 prospective-card checks (those files are not included), and re-running the low-rank comparison men-
 217 tioned in the supplement needs a checkout of BenchPress’s released code.

218 **Rows are not independent.** Up to sixteen benchmark rows come from a single quantization run.
 219 We mitigate with group-conditional calibration and the support floor, and note that a bootstrap over
 220 the 19 quantized checkpoints gives $[83.2, 94.7]\%$, close to the row-level interval, but the unit of anal-
 221 ysis remains a row. The leave-family-out coverage measurement has the same dependence: each
 222 row is scored once under each of 7 calibration families, so the 5719 evaluations are not indepen-
 223 dent, and every coverage judgment in this paper resamples whole checkpoints rather than counting
 224 evaluations.

225 9 Audit

226 We record errors found in our own work, including those that changed published claims. This
 227 section exists because the alternative—reporting only what survived—would misrepresent how the
 228 result was obtained.

229 Errors that changed a published number.

- 230 • *A retracted headline.* An early claim that held-out models came from families never used
 231 in construction was wrong; they were new checkpoints of families already present. The
 232 strict figure is 118/131.
- 233 • *Two parser bugs.* A benchmark label typo cost 11 training rows; separately, MATH-500 was
 234 matched as Math-Lv1-5, contaminating 7 test rows with scores from a different bench-
 235 mark.
- 236 • *A centring mismatch.* Calibration measured residuals about the stratum mean while in-
 237 tervals centred on the scheme mean, mis-sizing every calibrated width. Found only by
 238 independent reimplementations.
- 239 • *A column-order fault.* One card template orders columns [recovery, base, quant], which our
 240 parser turned into a spurious -68pp delta. The integrity gate rejected 10 of 11 rows on that
 241 card, passing only the one where base = quant made the check degenerate.
- 242 • *A retracted single-publisher claim.* We had asserted that one publisher was close to the
 243 only source of paired evaluations. A census of 10 publishers found 10 paired-evaluation
 244 cards outside the control publisher, 9 of them Intel’s; what differs is whose rows can be
 245 verified (§8).
- 246 • *A bootstrap on the wrong statistic.* The checkpoint-level bootstrap in our coverage audit re-
 247 sampled the two-sided indicator while the verdict was judged on the one-sided statistic. An
 248 external commenter asked for it to wrap the one-sided figure, which exposed the mismatch
 249 (§4).

250 **A regression we nearly shipped.** An audit indicated that a raw asymmetric empirical band cov-
 251 ered better than the shipped symmetric conformal interval (93.4% versus 89.6%). Since quantization
 252 damage is left-skewed, a symmetric band is obviously the wrong shape, and we were about to switch.
 253 On measuring first, the empirical band returns an *infinite* interval on 43% of rows, because a two-
 254 sided empirical index requires more calibration points than are often available. An infinite interval
 255 covers by construction. On rows where both are finite, the original wins on coverage *and* width. The
 256 change was reversed.

257 **An overclaimed guarantee.** The artifact was described as carrying a coverage guarantee without
258 qualification. Checking the method against the conformal literature rather than against our own
259 earlier work showed the guarantee is conditional on exchangeability that selection breaks. Corrected
260 in §8.

261 **Provenance discipline.** Two parser fixes were informed by inspecting rows that were already in
262 the test set. Rather than silently retain them, we record which fixes were test-informed and exclude
263 the affected family from the strict headline.

264 **Mechanical verification.** The figures this paper reports as results are macros generated from a
265 claims registry that recomputes each value from source artifacts, and the accompanying code fails a
266 test if an outward-facing document contains an untraced numeric literal. The exceptions are histor-
267 ical figures in the error list above (row counts and the -68pp delta) and the 90% marginal-interval
268 figure in §4, which are quoted from the provenance record rather than recomputed.

269 10 Conclusion

270 Per-model prediction of quantization accuracy damage does not work on published data, and the
271 reason is largely measurement noise rather than model capacity. What does work is narrower: a
272 per-scheme calibrated envelope, validated at 90.1% coverage on unseen checkpoints for 4 of the 6
273 schemes, and that says which of its own cells it cannot judge. Each cell is trusted; or insufficient
274 evidence, in which case the interval is still shown but flagged and the scheme demoted; or refused,
275 in which case the interval is withheld. Today 9 of 17 cells are insufficient evidence and 0 are refused.
276 The more useful contribution may be the limitation itself—the published record of quantized models
277 is outcome-truncated, and we measured, on deliberately faulted recipes for models up to 1.5B, losses
278 more than four times worse than anything it contains. A risk estimate built on that record is a floor
279 on risk, not a ceiling, and should say so.

280 References

- 281 Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. Conformal pre-
282 diction beyond exchangeability. *The Annals of Statistics*, 51(2), 2023. doi: 10.1214/23-AOS2276.
- 283 Henrik Boström and Ulf Johansson. Mondrian conformal regressors. In *Conformal and Probabilistic*
284 *Prediction and Applications*, volume 128 of *Proceedings of Machine Learning Research*, pages
285 1–20, 2020.
- 286 Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. GPTQ: Accurate post-training
287 quantization for generative pre-trained transformers. *arXiv preprint arXiv:2210.17323*, 2022.
- 288 Prashansa Gupta and Sean MacAvaney. On survivorship bias in MS MARCO. In *Proceedings*
289 *of the 45th International ACM SIGIR Conference on Research and Development in Information*
290 *Retrieval*, 2022. doi: 10.1145/3477495.3531832.
- 291 Ying Jin and Zhimei Ren. Confidence on the focal: Conformal prediction with selection-conditional
292 coverage. *arXiv preprint arXiv:2403.03868*, 2024.
- 293 Tanishq Kumar, Zachary Ankner, Benjamin F. Spector, Blake Bordelon, Niklas Muennighoff, Man-
294 sheej Paul, Cengiz Pehlevan, Christopher Ré, and Aditi Raghunathan. Scaling laws for precision,
295 2024.
- 296 Eldar Kurtić, Alexandre Marques, Mark Kurtz, and Dan Alistarh. We ran over half
297 a million evaluations on quantized LLMs—here’s what we found. Red Hat De-
298 veloper, October 2024. [https://developers.redhat.com/articles/2024/10/17/](https://developers.redhat.com/articles/2024/10/17/we-ran-over-half-million-evaluations-quantized-llms)
299 [we-ran-over-half-million-evaluations-quantized-llms](https://developers.redhat.com/articles/2024/10/17/we-ran-over-half-million-evaluations-quantized-llms).
- 300 Jing Lei, Max G’Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry Wasserman. Distribution-
301 free predictive inference for regression. *Journal of the American Statistical Association*, 113(523):
302 1094–1111, 2018. doi: 10.1080/01621459.2017.1307116.

- 303 Yujia Tong, Yuxi Wang, Yunyang Wan, Tian Zhang, Junhao Dong, and Jingling Yuan. Does com-
304 pression preserve uncertainty? a unified benchmark for quantized and sparse LLMs via conformal
305 prediction. *arXiv preprint arXiv:2606.01850*, 2026. Wuhan University of Technology; Nanyang
306 Technological University.
- 307 vLLM Project. llm-compressor, 2024. [https://github.com/vllm-project/
308 llm-compressor](https://github.com/vllm-project/llm-compressor).
- 309 Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random*
310 *World*. Springer, 2005. doi: 10.1007/b106715.
- 311 Guangxuan Xiao, Ji Lin, Mickael Seznec, Hao Wu, Julien Demouth, and Song Han. SmoothQuant:
312 Accurate and efficient post-training quantization for large language models. *arXiv preprint*
313 *arXiv:2211.10438*, 2022.
- 314 Yuchen Zeng and Dimitris Papailiopoulos. You don’t need to run every eval. *arXiv preprint*
315 *arXiv:2606.24020*, 2026. Microsoft Research, AI Frontiers.